

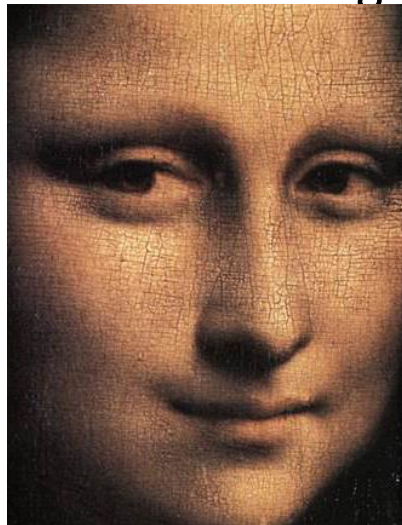
ECS797P - Machine Learning for Visual Data Analysis

Tao Xiang

t.xiang@qmul.ac.uk

Lecture 3

An instance classification
problem: Face Recognition



Face Recognition

- Face is the most common biometric used by humans
- Applications range from static, mug-shot verification to a dynamic, uncontrolled face identification in a cluttered background
- Challenges:
 - automatically locate the face
 - recognize the face from a general view point under different
 - illumination conditions
 - facial expressions
 - aging effects

Authentication vs. Identification

- Face Authentication/Verification (1:1 matching)



- Face Identification/Recognition (1:N matching)



Applications

□ Access Control



www.viisage.com



VISIONICS CORPORATION
Empowering Identification™

IBIS - Mobile Identification

Captures forensic quality fingerprints and photographs
Wirelessly processes and transmits data
Interfaces with all major databases (AFIS, NCIC 2000, etc.)
Provides real-time remote identification

www.visionics.com

www.visionics.com

Applications

□ Video Surveillance (On-line or off-line)

Face Scan at Airports



The St. Petersburg-Clearwater Airport installed facial recognition systems at two security checkpoints in January. Six-foot tall towers (above) house cameras that snap pictures of passengers as they pass through magnetometers. The passengers' faces instantly are compared to a database of images of wanted criminals. Sheriff Everett Rice (above left) was one of the first people to pass through the new security system.



www.facesnap.de

Why is Face Recognition Hard?

Many faces of Madonna



Face Recognition Difficulties

- Identify similar faces (**inter-class similarity**)
- Accommodate **intra-class variability** due to:
 - head pose
 - illumination conditions
 - expressions
 - facial accessories
 - aging effects

Inter-class Similarity

- Different persons may have very similar **appearance**



www.marykateandashley.com

Twins

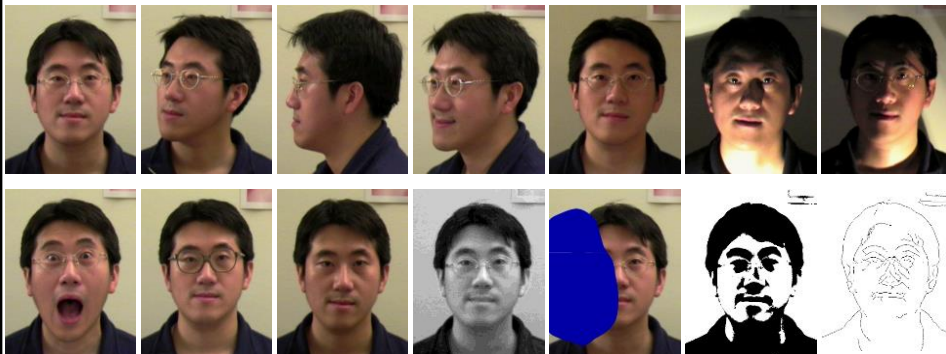


news.bbc.co.uk/1/english/in_depth/americas/2000/us_elections

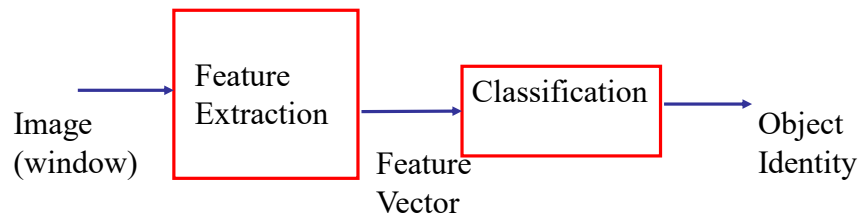
Father and son

Intra-class Variability

- Faces with intra-subject variations in pose, illumination, expression, accessories, color, occlusions, and brightness

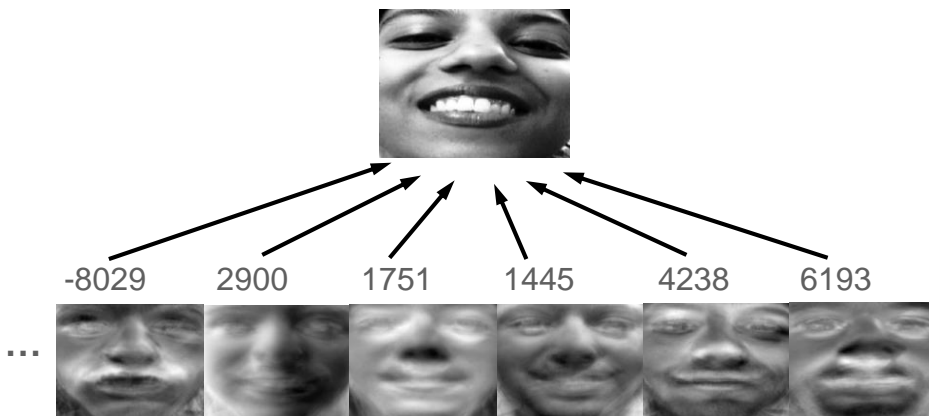


Sketch of a Pattern Recognition Architecture



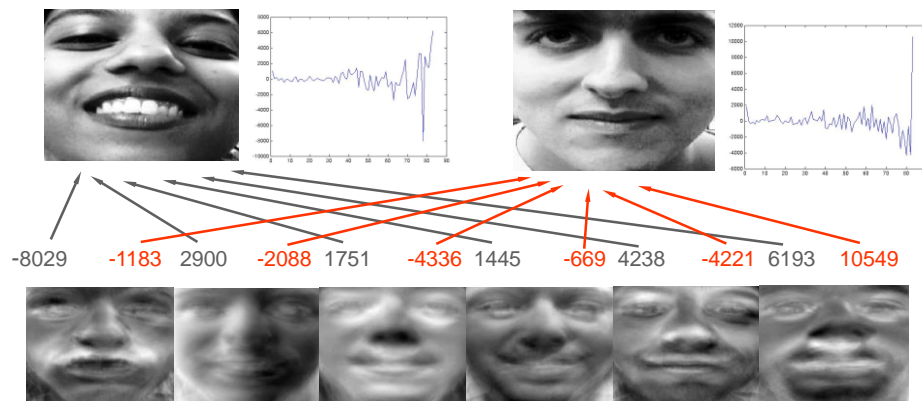
Eigenfaces: the idea

- Think of a face as being a weighted combination of some “component” or “basis” faces
- These basis faces are called eigenfaces



Eigenfaces: representing faces

- These basis faces can be differently weighted to represent any face
- So we can use different vectors of weights to represent different faces



Learning Eigenfaces

Q: How do we pick the set of basis faces?

A: We take a set of real training faces



Then we find (learn) a set of basis faces which best represent the differences between them

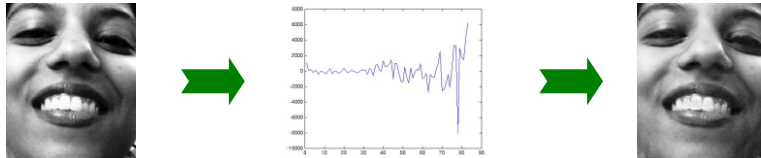
We'll use a statistical criterion for measuring this notion of "best representation of the differences between the training faces"

We can then store each face as a set of weights for those basis faces

Using Eigenfaces: recognition & reconstruction

- We can use the eigenfaces in two ways

1: we can store and then reconstruct a face from a set of weights



2: we can recognise a new picture of a familiar face

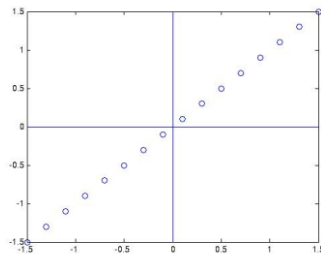


Learning Eigenfaces

- How do we learn them?
- We use a method called Principle Components Analysis (PCA)
- To understand this we will need to understand
 - What an eigenvector is
 - What covariance is
- But first we will look at what is happening in PCA qualitatively

Subspaces

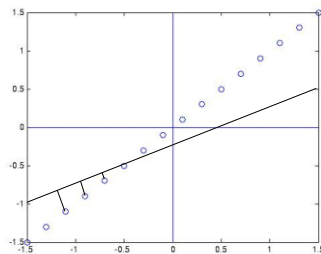
- Imagine that our face is simply a (high dimensional) vector of pixels
- We can think more easily about 2D vectors



- Here we have data in two dimensions
- But we only really need one dimension to represent it

Finding Subspaces

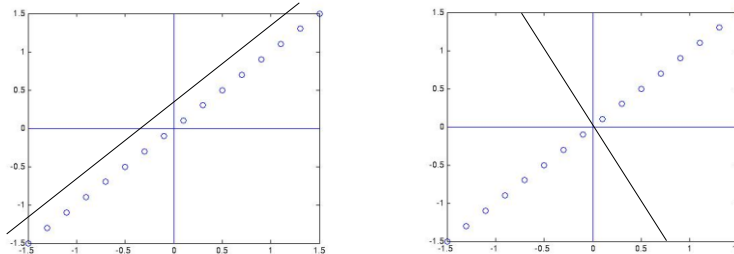
- Suppose we take a line through the space



- And then take the projection of each point onto that line
- This could represent our data in “one” dimension

Finding Subspaces

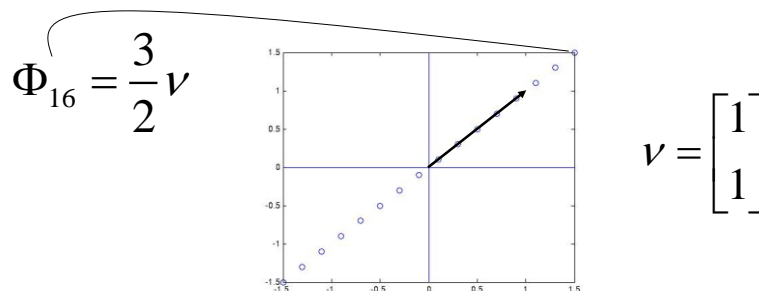
- Some lines will represent the data in this way well, some badly



- This is because the projection onto some lines separates the data well, and the projection onto some lines separates it badly

Finding Subspaces

- Rather than a line we can perform roughly the same trick with a vector ν



- Now we have to scale the vector to obtain any point on the line

$$\Phi_i = \mu \nu$$

Eigenvectors

- An eigenvector is a vector v that obeys the following rule:

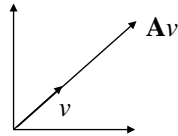
$$Av = \mu v$$

Where A is a matrix, μ is a scalar (called the eigenvalue)

e.g. $A = \begin{bmatrix} 2 & 3 \\ 2 & 1 \end{bmatrix}$ one eigenvector of A is $v = \begin{bmatrix} 3 \\ 2 \end{bmatrix}$ since

$$\begin{bmatrix} 2 & 3 \\ 2 & 1 \end{bmatrix} \begin{bmatrix} 3 \\ 2 \end{bmatrix} = \begin{bmatrix} 12 \\ 8 \end{bmatrix} = 4 \times \begin{bmatrix} 3 \\ 2 \end{bmatrix}$$

so for this eigenvector of this matrix the eigenvalue is 4

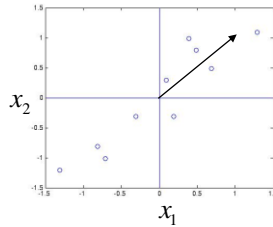


Eigenvectors

- We can think of matrices as performing transformations on vectors (e.g rotations, reflections)
- We can think of the eigenvectors of a matrix as being special vectors (for that matrix) that are scaled by that matrix
- Different matrices have different eigenvectors
- Only square matrices have eigenvectors
- Not all square matrices have eigenvectors
- An n by n matrix has at most n distinct eigenvectors
- All the distinct eigenvectors of a matrix are orthogonal (i.e., perpendicular)

Covariance

- Which single vector can be used to separate these points as much as possible?



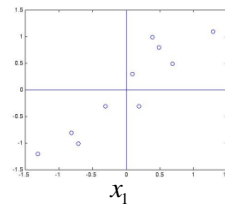
- This vector turns out to be a vector expressing the direction of the correlation
- Here I have two variables x_1 and x_2
- They co-vary (y tends to change in roughly the same direction as x)

Covariance

- The covariances can be expressed as a matrix

$$C = \begin{bmatrix} x_1 & x_2 \\ .617 & .615 \\ .615 & .717 \end{bmatrix} \begin{matrix} x_1 \\ x_2 \end{matrix}$$

The matrix C is shown with its elements. The first row contains the variable names x_1 and x_2 . The first column contains the values $.617$ and $.615$, and the second column contains $.615$ and $.717$. The matrix is labeled C on the left. A curved arrow points from the matrix to the scatter plot on the right.



- The diagonal elements are the variances e.g. $\text{Var}(x_1)$
- The covariance of two variables is:

$$\text{cov}(x_1, x_2) = \frac{\sum_{i=1}^n (x_1^i - \bar{x}_1)(x_2^i - \bar{x}_2)}{n-1}$$

Eigenvectors of the covariance matrix

- The covariance matrix has eigenvectors

covariance matrix

$$C = \begin{bmatrix} .617 & .615 \\ .615 & .717 \end{bmatrix}$$

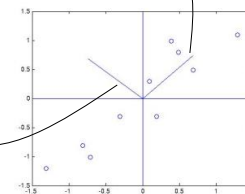
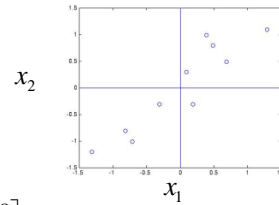
eigenvectors

$$v_1 = \begin{bmatrix} -.735 \\ .678 \end{bmatrix} \quad v_2 = \begin{bmatrix} .678 \\ .735 \end{bmatrix}$$

eigenvalues

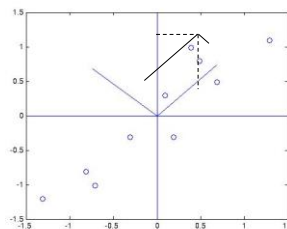
$$\mu_1 = 0.049 \quad \mu_2 = 1.284$$

- Eigenvectors with larger eigenvalues correspond to directions in which the data varies more
- Finding the eigenvectors and eigenvalues of the covariance matrix for a set of data is termed principle components analysis



Expressing points using eigenvectors

- Suppose you think of your eigenvectors as specifying a new vector space



- i.e. I can reference any point in terms of those eigenvectors
- A point's position in this new coordinate system is what we earlier referred to as its "weight vector"
- For many data sets you can cope with fewer dimensions in the new space than in the old space

Eigenfaces

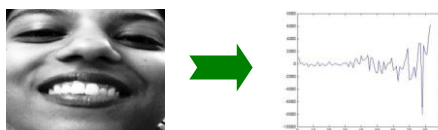
- All we are doing in the face case is treating the face as a point in a high-dimensional space, and then treating the training set of face pictures as our set of points
- To train:
 - We calculate the covariance matrix of the faces
 - We then find the eigenvectors of that covariance matrix
- These eigenvectors are the eigenfaces or basis faces
- Eigenfaces with bigger eigenvalues will explain more of the variation in the set of faces, i.e. will be more distinguishing

Eigenfaces: image space to face space

- When we see an image of a face we can transform it to face space

$$\mathbf{w}_k = \mathbf{x}^i \cdot \mathbf{v}_k$$

- There are $k=1\dots n$ eigenfaces \mathbf{v}_k
- The i^{th} face in image space is a vector \mathbf{x}^i
- The corresponding weight is \mathbf{w}_k
- We calculate the corresponding weight for every eigenface

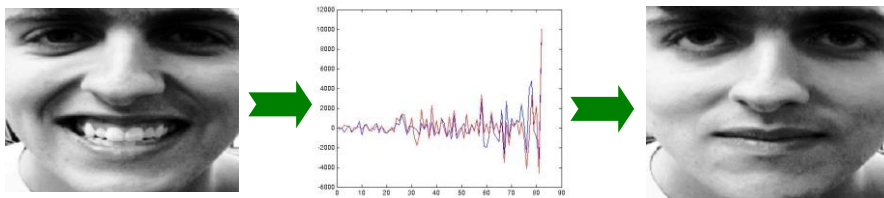


Recognition in face space

- Recognition is now simple. We find the euclidean distance d between our face and all the other stored faces in face space:

$$d(w^1, w^2) = \sqrt{\sum_{i=1}^n (w_i^1 - w_i^2)^2}$$

- The closest face in face space is the chosen match



Reconstruction

- The more eigenfaces you have the better the reconstruction, but you can have high quality reconstruction even with a small number of eigenfaces

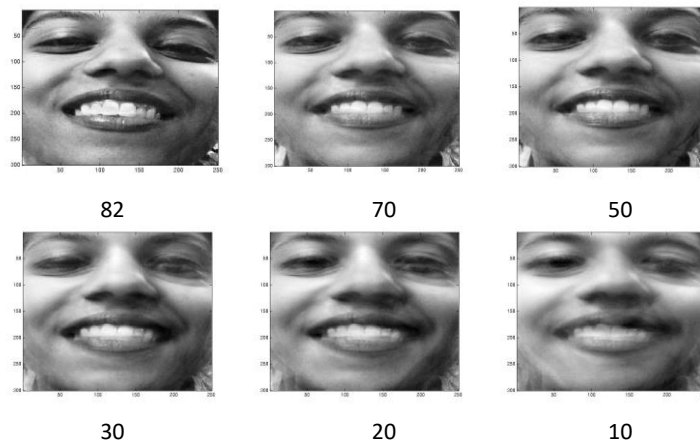


Image Representation

- Training set of m images of size $N \times N$ are represented by vectors of size N^2

$$x_1, x_2, x_3, \dots, x_m$$

Example

$$\begin{bmatrix} 1 & 2 & 3 \\ 3 & 1 & 2 \\ 4 & 5 & 1 \end{bmatrix}_{3 \times 3} \longrightarrow \begin{bmatrix} 1 \\ 2 \\ 3 \\ 3 \\ 1 \\ 2 \\ 4 \\ 5 \\ 1 \end{bmatrix}_{9 \times 1}$$



Average Image and Difference Images

- The average training set is defined by

$$\mu = (1/m) \sum_{i=1}^m x_i$$



- Each face differs from the average by vector

$$r_i = x_i - \mu$$

Covariance Matrix

- The covariance matrix is constructed as

$$C = AA^T \text{ where } A=[r_1, \dots, r_m]$$

Size of this matrix is $N^2 \times N^2$

- Finding eigenvectors of $N^2 \times N^2$ matrix is intractable. Hence, use the matrix $A^T A$ of size $m \times m$ and find eigenvectors of this small matrix.

Eigenvalues and Eigenvectors - Definition

- If v is a nonzero vector and λ is a number such that

$$Av = \lambda v, \text{ then}$$

v is said to be an *eigenvector* of A with *eigenvalue* λ .

Example

$$\begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix} \times \begin{bmatrix} 1 \\ 1 \end{bmatrix} = 3 \times \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

λ (eigenvalues)

A

V (eigenvectors)

Eigenvectors of Covariance Matrix

- The eigenvectors v_i of $A^T A$ are:

$$\begin{matrix} \begin{bmatrix} 0.3923 \\ 0.5060 \\ 0.7681 \end{bmatrix} & \begin{bmatrix} 0.9087 \\ -0.0835 \\ -0.4091 \end{bmatrix} & \begin{bmatrix} 0.1429 \\ -0.8585 \\ 0.4926 \end{bmatrix} \\ v_1 & v_2 & v_3 \end{matrix}$$

- Consider the eigenvectors v_i of $A^T A$ such that $A^T A v_i = \mu_i v_i$

- Premultiplying both sides by A , we have $AA^T(Av_i) = \mu_i(Av_i)$

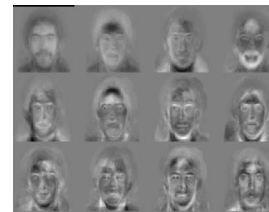
Face Space

- The eigenvectors of covariance matrix are

$$u_i = A v_i$$

$$u_i = A v_i \quad \text{Face Space}$$

$$U = \begin{matrix} \begin{bmatrix} -0.2621 \\ -0.2621 \\ -0.6527 \\ -0.1137 \\ -0.5589 \\ -0.6015 \\ -0.4895 \\ 0 \\ 0.6379 \end{bmatrix} & \begin{bmatrix} 0.3256 \\ 0.3256 \\ -3.3773 \\ 0.9922 \\ -1.0076 \\ -0.5080 \\ 2.3100 \\ -1.6434 \\ 1.7008 \end{bmatrix} & \begin{bmatrix} -1.3511 \\ -1.3511 \\ 2.2735 \\ 1.0014 \\ -6.0561 \\ -5.4206 \\ 0.6517 \\ 0 \\ 1.7008 \end{bmatrix} \\ u_1 & u_2 & u_3 \end{matrix}$$



- u_i resemble facial images which look ghostly, hence called **Eigenfaces**

Projection into Face Space

- A face image can be projected into this face space by

$$p_k = U^T(x_k - \mu) \text{ where } k=1,\dots,m$$

Recognition

- The test image x is projected into the face space to obtain a vector p :

$$p = U^T(x - \mu)$$

- The distance of p to each face class is defined by

$$\epsilon_k^2 = ||p - p_k||^2; k = 1,\dots,m$$

- A distance threshold Θ_c is half the largest distance between any two face images:

$$\Theta_c = \frac{1}{2} \max_{j,k} \{ ||p_j - p_k|| \}; j,k = 1,\dots,m$$

Recognition

- Find the distance ϵ between the original image x and its reconstructed image from the eigenface space, x_f ,
$$\epsilon^2 = ||x - x_f||^2, \text{ where } x_f = U * p + \mu$$
- Recognition process:
 - IF $\epsilon \geq \theta_c$
then input image is not a face image;
 - IF $\epsilon < \theta_c$ AND $\epsilon_k \geq \theta_c$ for all k
then input image contains an unknown face;
 - IF $\epsilon < \theta_c$ AND $\epsilon_k^* = \min_k \{ \epsilon_k \} < \theta_c$
then input image contains the face of individual k^*

Limitations of Eigenfaces Approach

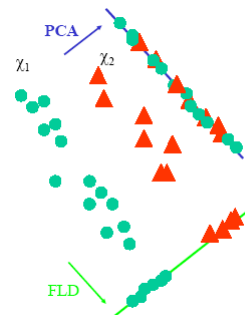
- **Variations in lighting conditions**
 - Different lighting conditions for enrolment and query.
 - Bright light causing image saturation.



- **Differences in pose – Head orientation**
 - 2D feature distances appear to distort.
- **Expression**
 - Change in feature location and shape.

Linear Discriminant Analysis

- PCA does not use class information
 - PCA projections are optimal for reconstruction from a low dimensional basis, they may not be optimal from a discrimination standpoint.
- LDA is an enhancement to PCA
 - Constructs a discriminant subspace that minimizes the scatter between images of same class and maximizes the scatter between different class images



Mean Images

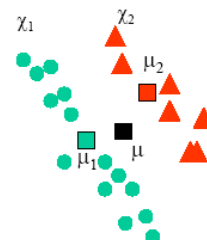
- Let X_1, X_2, \dots, X_c be the face classes in the database and let each face class X_i , $i = 1, 2, \dots, c$ has k facial images x_j , $j = 1, 2, \dots, k$.

- We compute the mean image μ_i of each class X_i as:

$$\mu_i = \frac{1}{k} \sum_{j=1}^k x_j$$

- Now, the mean image μ of all the classes in the database can be calculated as:

$$\mu = \frac{1}{c} \sum_{i=1}^c \mu_i$$



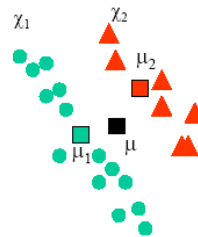
Scatter Matrices

- We calculate **within-class** scatter matrix as:

$$S_W = \sum_{i=1}^c \sum_{x_k \in X_i} (x_k - \mu_i)(x_k - \mu_i)^T$$

- We calculate the **between-class** scatter matrix as:

$$S_B = \sum_{i=1}^c N_i (\mu_i - \mu)(\mu_i - \mu)^T$$



Multiple Discriminant Analysis

We find the projection directions as the matrix W that maximizes

$$\hat{W} = \operatorname{argmax}_W J(W) = \frac{|W^T S_B W|}{|W^T S_W W|}$$

This is a generalized Eigenvalue problem where the columns of W are given by the vectors w_i that solve

$$S_B w_i = \lambda_i S_W w_i$$

Fisherface Projection

- We find the product of S_W^{-1} and S_B and then compute the Eigenvectors of this product ($S_W^{-1} S_B$) - AFTER REDUCING THE DIMENSION OF THE FEATURE SPACE.
- Use same technique as Eigenfaces approach to reduce the dimensionality of scatter matrix to compute eigenvectors.
- Form a matrix W that represents all eigenvectors of $S_W^{-1} S_B$ by placing each eigenvector w_i as a column in W .
- Each face image $x_j \in X_i$ can be projected into this face space by the operation
$$p_i = W^T(x_j - \mu)$$



Testing

- Same as Eigenfaces Approach

References

- Turk, M., Pentland, A.: *Eigenfaces for recognition*. J. Cognitive Neuroscience **3** (1991) 71–86.
- Belhumeur, P., Hespanha, J., Kriegman, D.: *Eigenfaces vs. Fisherfaces: recognition using class specific linear projection*. IEEE Transactions on Pattern Analysis and Machine Intelligence **19** (1997) 711–720.

State of the art

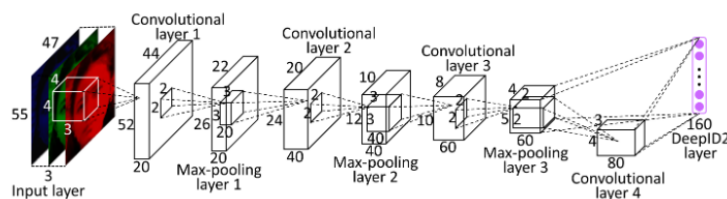


Figure 1: The ConvNet structure for DeepID2 extraction.

- Prof. X. Tang's group: DeepID2
<http://arxiv.org/pdf/1406.4773v1.pdf>
- Updated version: DeepID3
<http://arxiv.org/pdf/1502.00873.pdf>